

Value Iteration Algorithm

The Bellman equation is the basis of the value iteration algorithm for solving MDPs.

Let $U_t(I)$ be the utility value for state s at the t 'th iteration. The iteration step, called a Bellman update, looks like this:

$$U_{t+1}(I) = \max_A \left[R(I, A) + \sum_J P(J | I, A) \cdot U_t(J) \right]$$

Where the Rewards $R(I, A)$ for each state is the expected reward of taking action A in State I . That is :-

$$R(I, A) = \sum_J P(J | I, A) \cdot R(J, A, I)$$

The value iteration algorithm is as follows :-

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$, rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

TASK 1

Step Costs : [-5,-5,-5]

Gamma : 0.99

Delta : 0.001

Iterations : 112

From the task_1_trace.txt file obtained, we can make the following inferences about the policy :

- Whenever Lero has 0 stamina, the only optimal policy for that state is to RECHARGE
- Whenever Lero has 0 arrows, he can RECHARGE or DODGE but cannot SHOOT
- Most of the times, Lero is Risk Averse. That is he prefers to RECHARGE than to SHOOT in cases when his stamina = 50 instead of losing an arrow.

TASK 2

PART 1

Step Costs : [-0.25,-2.5,-2.5]

Gamma : 0.99

Delta : 0.001

Iterations : 100

From the task_2_part_1_trace.txt file obtained we can make the following inferences about the changes in the policy :

- Since the step cost for the SHOOT action is less negative compared to the other actions, Lero now becomes Risk Seeking and shoots whenever possible.
- Convergence is faster. In only 100 iterations, the value iteration algorithm converges when compared to the 112 iterations in TASK 1
- Magnitude of the utilities are less negative in this case compared to TASK 1 because of the less negative step cost.

PART 2

Step Costs : [-2.5,-2.5,-2.5]

Gamma : 0.1

Delta : 0.001

Iterations : 5

From the task_2_part_2_trace.txt file obtained we can make the following inferences about the policy :

- All states converge with a policy in the final iteration with more or less same utilities.
- The RECHARGE policy is preferred over the others hence indicating a Risk Averse behaviour.
- Bellman update is a contraction by a factor of γ on the space of utility vectors. There is fast convergence if we make γ small, but this effectively gives the agent a short horizon and could miss the long-term effects of the agent's actions. So we can see that, Lero tries to play it safe and keeps RECHARGING instead of looking into the future.

- Due to a very small discount factor (Gamma), the bellman iterations converge fast, within 4 iterations.

PART 3

Step Costs : [-2.5,-2.5,-2.5]

Gamma : 0.1

Delta : 0.0000000001

Iterations : 12

From the task_2_part_3_trace.txt file obtained we can make the following inferences about the policy :

- The only difference between part 2 and part 3 is delta.
- Since the Bellman error is very small we would require more iterations to get a higher accuracy. Clearly the number of iterations in part 3 are greater than the number of iterations in part 2.
- Lero appears to be risk neutral according to the policy obtained.